



Energy-Efficient Reconfigurable Systolic Array with Huffman Compression and Clock Gating

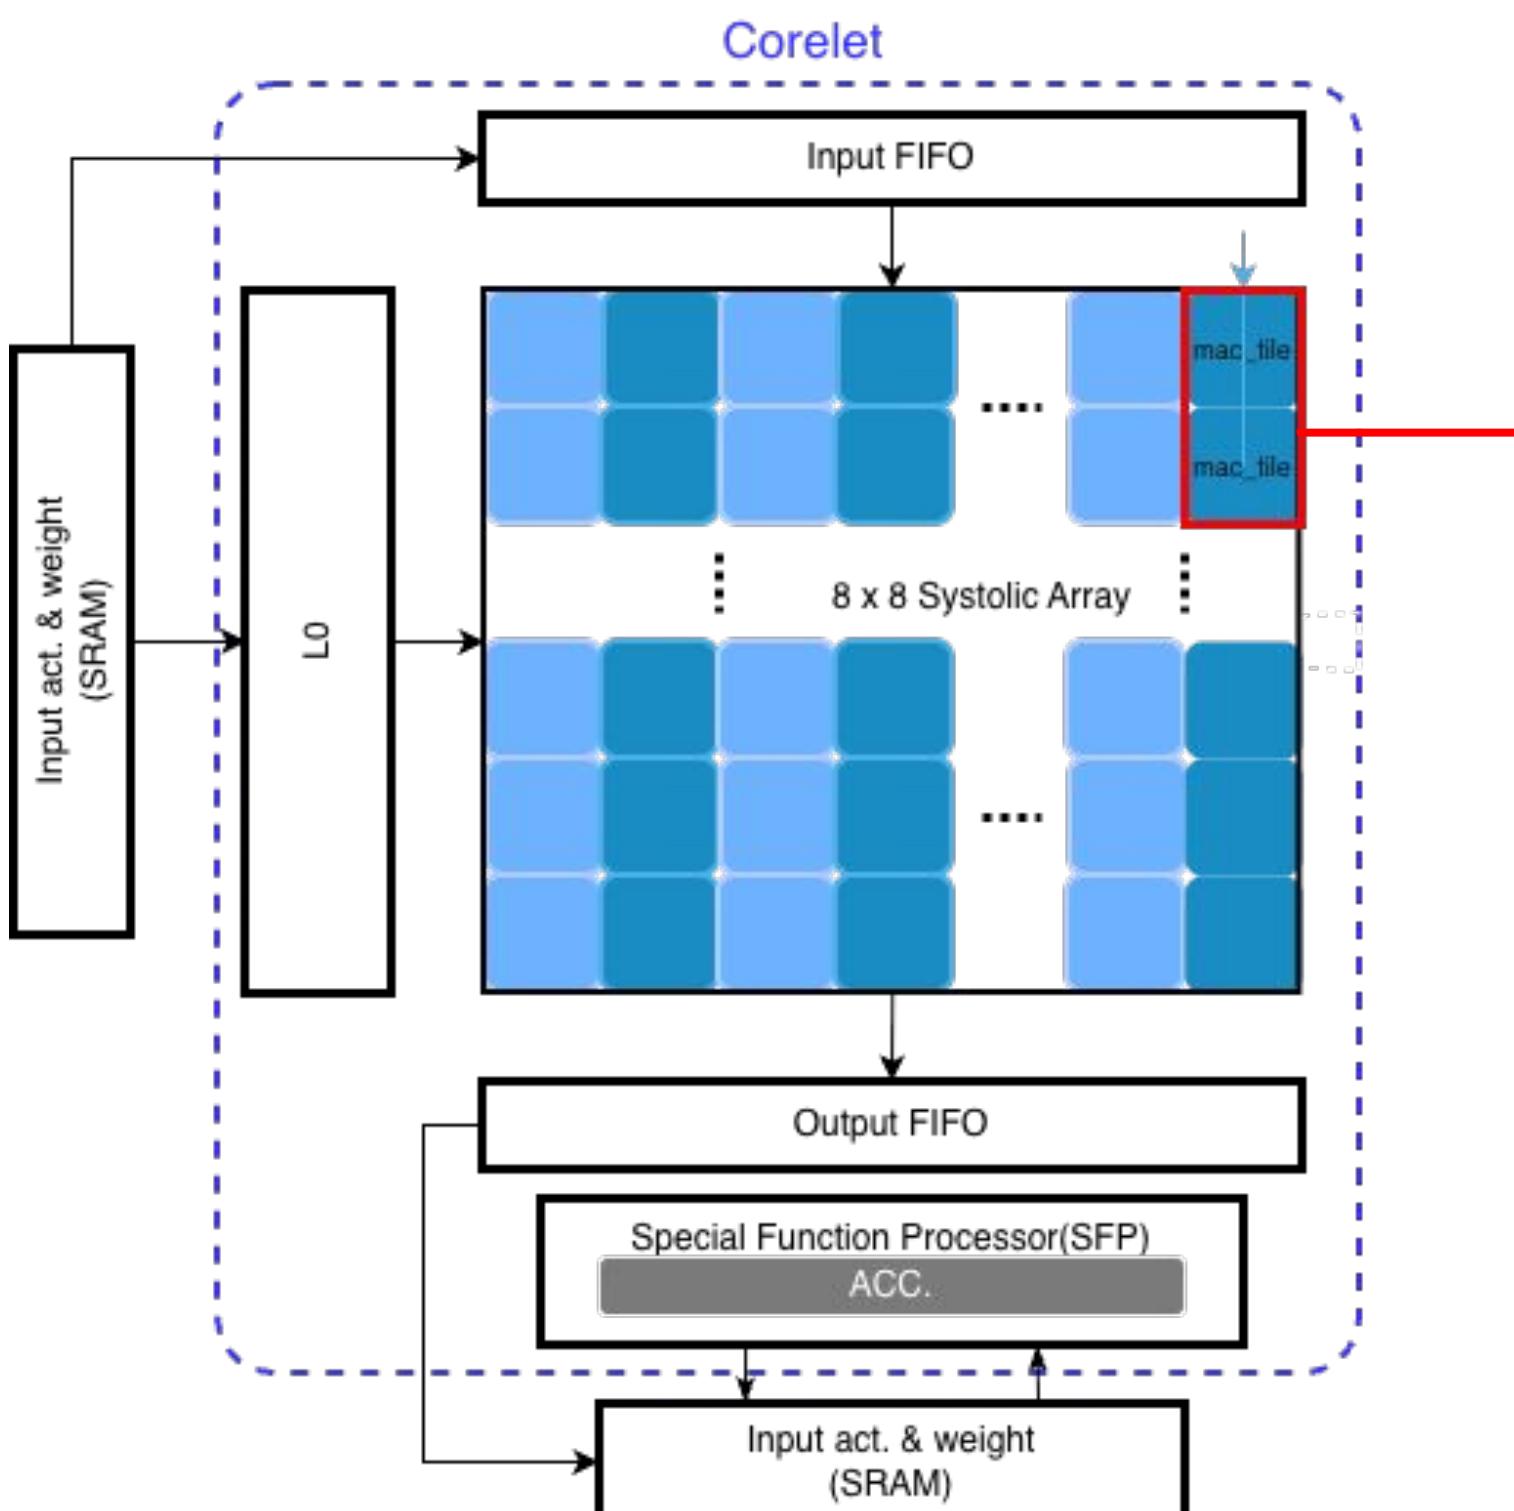


VLSI Scream

Chi-Han Chiu, Gary(Kai-Jui) Weng, Yun-Chen Tsai, Bing-Cheng Chiang, Yufan Wang

Foundation

- Combined Reconfigurability (Alpha 1)



| | ws_2bit | ws_4bit | os_2bit | os_4bit |
|------|-------------|---------|----------------|------------|
| in_n | psum | psum | {8'b0, w1, w0} | {12'b0, w} |
| in_w | act, w0, w1 | act, w | act | act |

Implemented on Cyclone 10 GX

| | Vanilla | Mixed |
|----------------|------------|-----------|
| Frequency | 125 MHz | 125 MHz |
| Total Register | 12,155 | 12,667 |
| Total DSP | 60 | 64 |
| LUT | 8,424 | 9,062 |
| Data Delay | 7.571 ns | 7.205 ns |
| Power | 1119.85 mW | 1256.2 mW |

Alpha 2 - Clock Gating

- Less power consumption than no clock gating design.

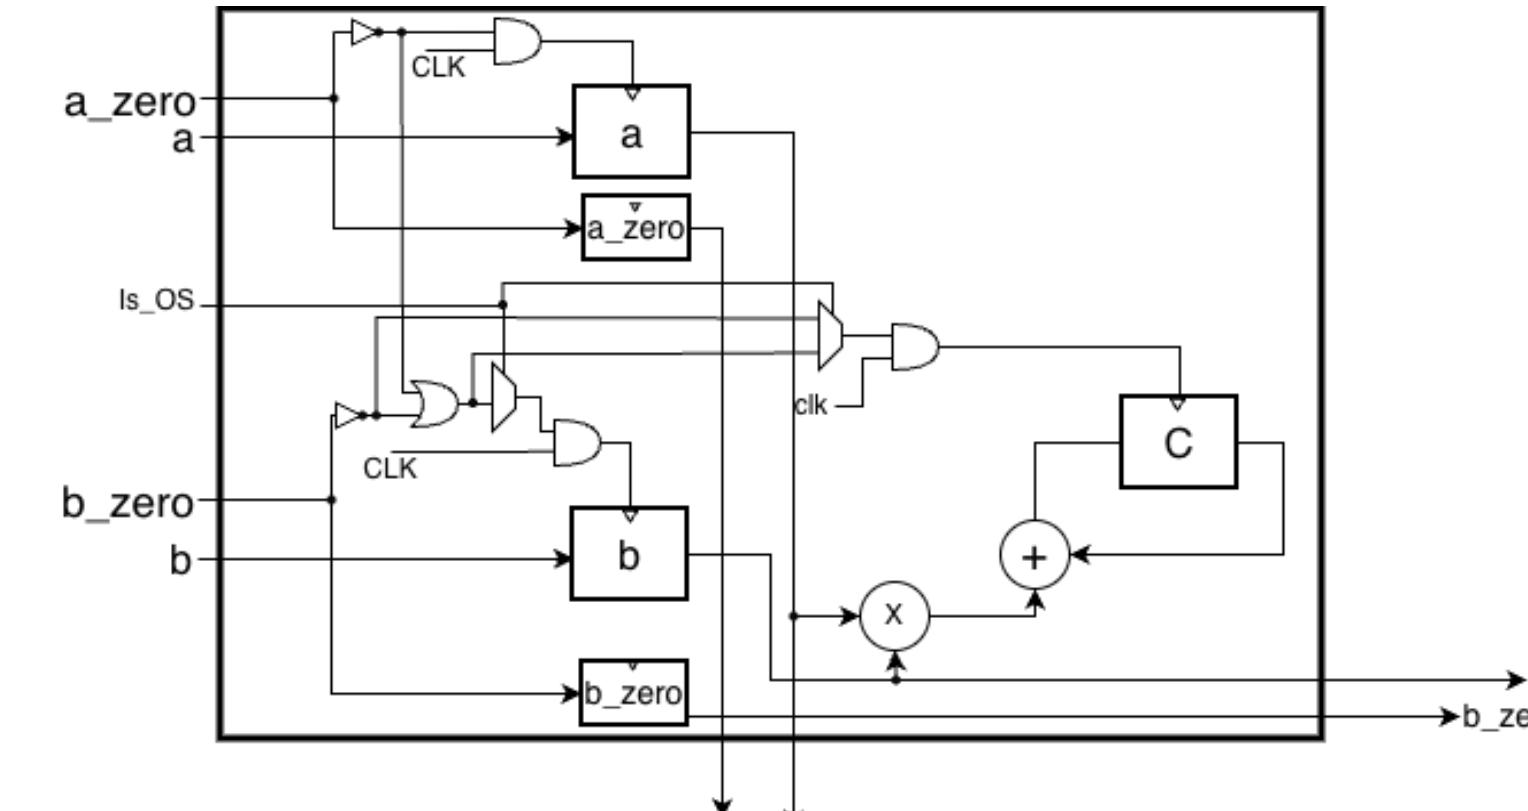


Table 1. Sparsity of Activation and Weight

$$\begin{aligned} \text{Sparsity}_{\text{new}} &= \text{Sparsity}_{\text{Act.}} + \text{Sparsity}_{\text{Wgt.}} \\ &\quad - \text{Sparsity}_{\text{Act.}} \times \text{Sparsity}_{\text{Wgt.}} \\ &= 86\% + 16\% - 86\% \times 16\% \\ &= 88.24\% \end{aligned}$$

- Our design can reach maximum of **88.24% power reduction**. (may be various to different FPGA.)
- To maximize the performance of this technique, the next target will be to enable fully shut down of the transistors.

Alpha 3 - Huffman Compression

| Value | 0000 | 0001 | 0010 | 0011 | 0100 | 0101 | 0110 | 0111 |
|-------------|------|------|------|------|------|------|------|------|
| Count | 248 | 11 | 10 | 5 | 4 | 4 | 4 | 2 |
| Save (bits) | -744 | -22 | -10 | 0 | +4 | +8 | +12 | +10 |

$$4 \times 288 - \sum \text{Save} = 410$$

- With Huffman Compression, we can get a **36.11% bit reduction** for activation.

Alpha 4 - Whole Conv. Layer

$$y = \frac{x - \mathbb{E}[x]}{\sqrt{\text{Var}[x] + \epsilon}} * \gamma + \beta$$

- Implementing **Batch Normalization, ReLU, and Max Pooling Layer** into our proposed 2-D systolic array.
- With these layers, it is able to fully go through the whole VGG16 CNN.

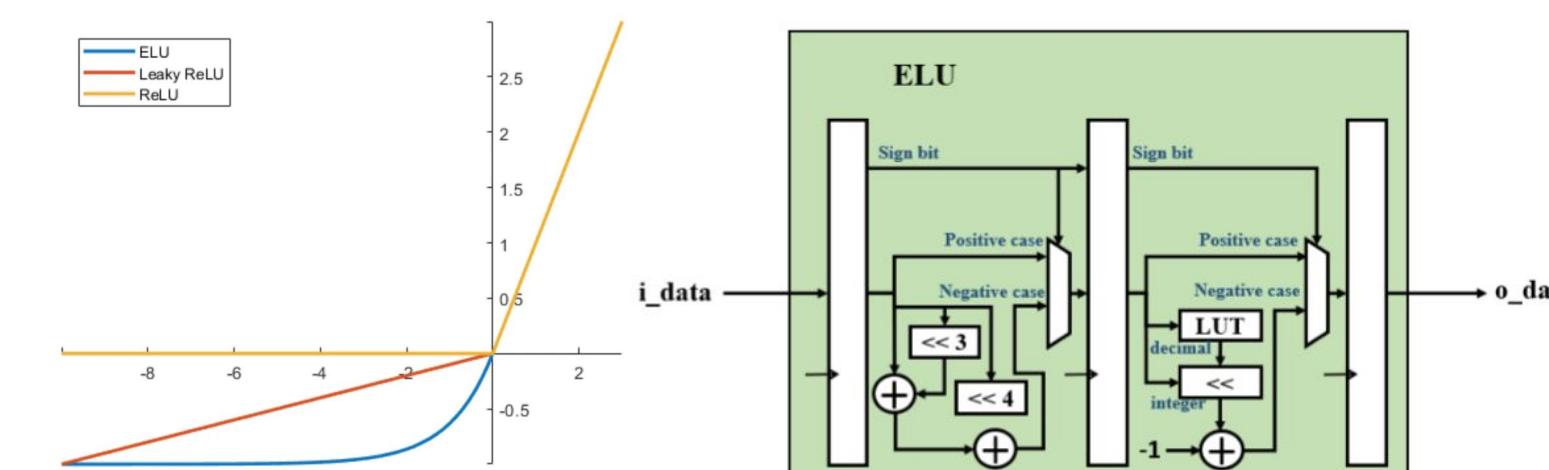
Alpha 5 - Model Fusion

$$\begin{cases} x_{conv} = w_{conv} * x_{in} + b \\ x_{bn} = \frac{(x_{conv} - \mu)}{\sqrt{\sigma^2 + \epsilon}} \gamma + \beta \end{cases} \quad \begin{cases} w' = \frac{\gamma}{\sqrt{\sigma^2 + \epsilon}} w_{conv} \\ b' = \frac{b - \mu}{\sqrt{\sigma^2 + \epsilon}} \gamma + \beta \end{cases}$$

$$x_{bn} = w' * x_{in} + b'$$

- BN parameters are fixed at inference and fused into adjacent Convolution layers, eliminating runtime normalization.
- This **reduces computation by 4.3% and hardware cost by 25%**, while simplifying datapath design and improving efficiency.

Alpha 6 - Flex Act. Func.

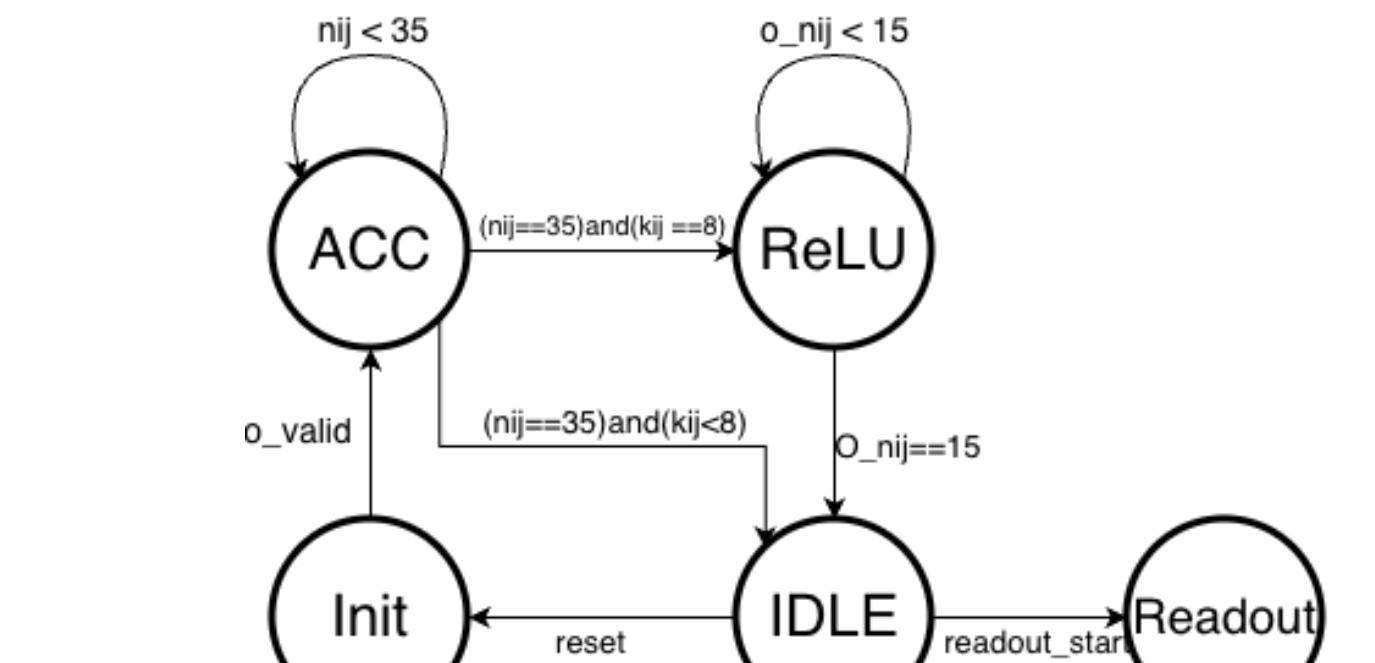


- Includes a flexible interface that supports all three of the most widely adopted activation functions—**ReLU, Leaky ReLU, and ELU**.

$$\exp(x) = 2^{\log_2 \exp(x)} = 2^{\frac{x}{\ln 2}}$$

- The right module is our implemented ELU Function. Since e^x is not hardware-friendly and its direct implementation incurs excessive area overhead, it is transformed into the above equivalent form.

Alpha 7 - FSM in SFU

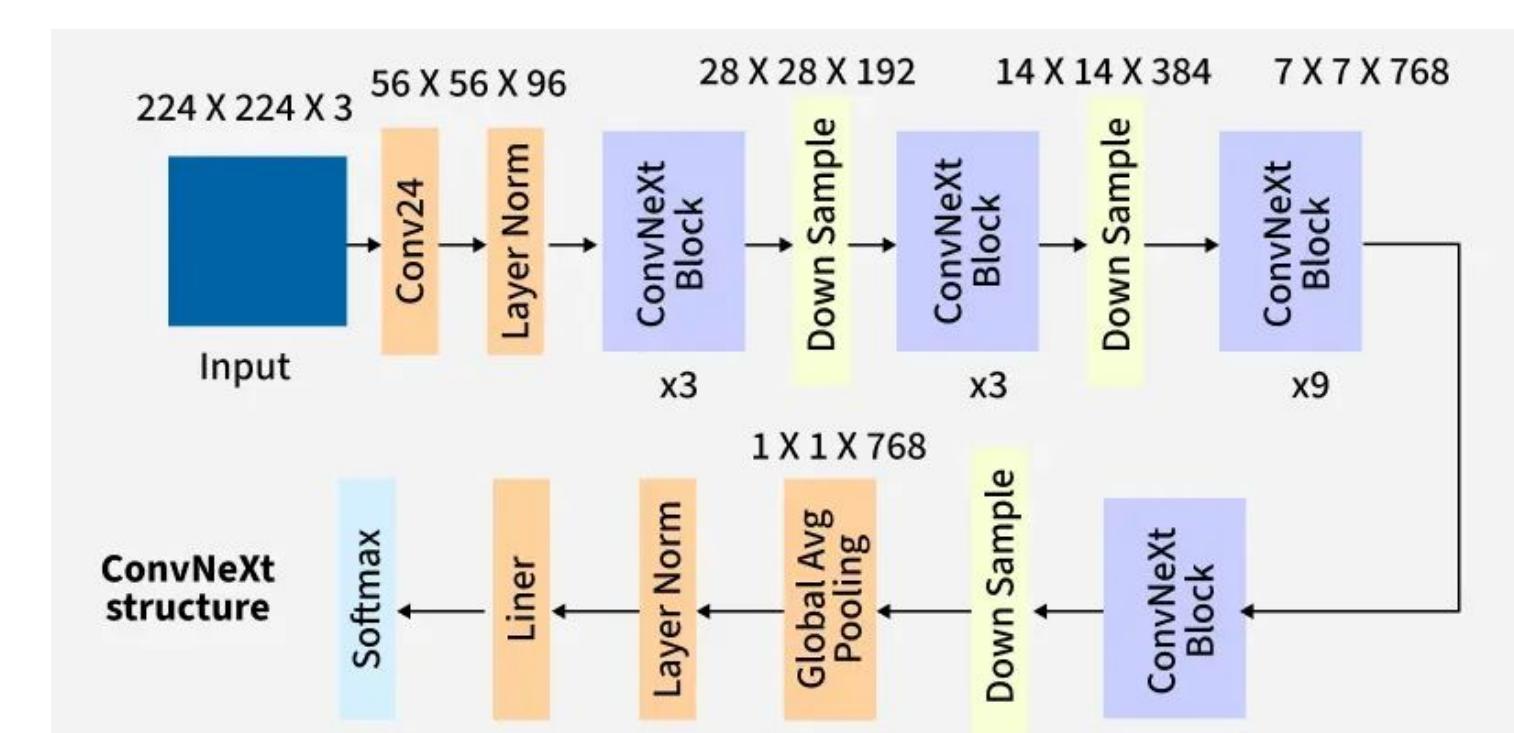


- TB has no control to PSUM memory.
- Inst_w is pipelined appropriately in core.v.

```
// high level instructions from TB
input [1:0] inst_w,
// L0 mem ctrls from TB
input CEN_xmem,
input WEN_xmem,
input [10:0] A_xmem,
input [bw*col-1:0] D_xmem, //32b

// from/to TB
input [3:0] kij, //for SFU
input readout_start, //trigger for output stage
output [psum_bw*col-1:0] readout //16*8b
```

Alpha 8 - ConvNext Application



- We aim to implement and evaluate ConvNeXt, a modern architecture known for achieving the same accuracy (~90%, 4-bit) with significantly fewer parameters.

| | VGGNet | ConvNext |
|------------|---------|----------|
| model size | 68.7 MB | 33.0 MB |